

Import Libraries & Load Data

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import joblib
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import confusion_matrix, classification_report,
accuracy_score, roc_auc_score, roc_curve
import warnings
warnings.filterwarnings('ignore')

# Load data
df = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv")
df.head()

    customerID  gender  SeniorCitizen Partner Dependents  tenure
PhoneService \
0  7590-VHVEG  Female           0      Yes        No       1
No
1  5575-GNVDE   Male            0      No        No      34
Yes
2  3668-QPYBK   Male            0      No        No       2
Yes
3  7795-CFOCW   Male            0      No        No      45
No
4  9237-HQITU  Female           0      No        No       2
Yes

    MultipleLines InternetService OnlineSecurity ...
DeviceProtection \
0  No phone service           DSL        No ...
No
1                  No          DSL        Yes ...
Yes
2                  No          DSL        Yes ...
No
3  No phone service           DSL        Yes ...
Yes
4                  No  Fiber optic        No ...
No

    TechSupport StreamingTV StreamingMovies          Contract
PaperlessBilling \
0          No         No        No Month-to-month
Yes
1          No         No        No One year
```

No				
2	No	No	No	Month-to-month
Yes				
3	Yes	No	No	One year
No				
4	No	No	No	Month-to-month
Yes				

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

Data Cleaning

```
# Check info
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customerID      7043 non-null   object 
 1   gender          7043 non-null   object 
 2   SeniorCitizen   7043 non-null   int64  
 3   Partner         7043 non-null   object 
 4   Dependents     7043 non-null   object 
 5   tenure          7043 non-null   int64  
 6   PhoneService    7043 non-null   object 
 7   MultipleLines   7043 non-null   object 
 8   InternetService 7043 non-null   object 
 9   OnlineSecurity  7043 non-null   object 
 10  OnlineBackup    7043 non-null   object 
 11  DeviceProtection 7043 non-null   object 
 12  TechSupport    7043 non-null   object 
 13  StreamingTV    7043 non-null   object 
 14  StreamingMovies 7043 non-null   object 
 15  Contract        7043 non-null   object 
 16  PaperlessBilling 7043 non-null   object 
 17  PaymentMethod   7043 non-null   object 
 18  MonthlyCharges  7043 non-null   float64
 19  TotalCharges    7043 non-null   object 
 20  Churn           7043 non-null   object 
```

```

dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

# Convert TotalCharges to numeric
df['TotalCharges']=pd.to_numeric(df['TotalCharges'],errors='coerce')
'''errors : {'ignore', 'raise', 'coerce'}, default 'raise'
   - If 'raise', then invalid parsing will raise an exception.
   - If 'coerce', then invalid parsing will be set as NaN.
   - If 'ignore', then invalid parsing will return the input'''

"errors : {'ignore', 'raise', 'coerce'}, default 'raise'\n      - If
'raise', then invalid parsing will raise an exception.\n      - If
'coerce', then invalid parsing will be set as NaN.\n      - If 'ignore',
then invalid parsing will return the input"

#check null value
df.isnull().sum()

customerID          0
gender              0
SeniorCitizen       0
Partner             0
Dependents          0
tenure              0
PhoneService        0
MultipleLines        0
InternetService     0
OnlineSecurity      0
OnlineBackup         0
DeviceProtection    0
TechSupport          0
StreamingTV          0
StreamingMovies      0
Contract            0
PaperlessBilling     0
PaymentMethod        0
MonthlyCharges       0
TotalCharges         11
Churn               0
dtype: int64

#fill null value with median
df['TotalCharges']=df['TotalCharges'].fillna(df['TotalCharges'].median())

df.head(20)

  customerID  gender  SeniorCitizen Partner Dependents  tenure
PhoneService \
0  7590-VHVEG  Female           0      Yes        No       1
No

```

1	5575-GNVDE	Male	0	No	No	34
Yes						
2	3668-QPYBK	Male	0	No	No	2
Yes						
3	7795-CFOCW	Male	0	No	No	45
No						
4	9237-HQITU	Female	0	No	No	2
Yes						
5	9305-CDSKC	Female	0	No	No	8
Yes						
6	1452-KI0VK	Male	0	No	Yes	22
Yes						
7	6713-OKOMC	Female	0	No	No	10
No						
8	7892-P00KP	Female	0	Yes	No	28
Yes						
9	6388-TABGU	Male	0	No	Yes	62
Yes						
10	9763-GRSKD	Male	0	Yes	Yes	13
Yes						
11	7469-LKBCI	Male	0	No	No	16
Yes						
12	8091-TTVAX	Male	0	Yes	No	58
Yes						
13	0280-XJGEX	Male	0	No	No	49
Yes						
14	5129-JLPIS	Male	0	No	No	25
Yes						
15	3655-SNQYZ	Female	0	Yes	Yes	69
Yes						
16	8191-XWSZG	Female	0	No	No	52
Yes						
17	9959-W0FKT	Male	0	No	Yes	71
Yes						
18	4190-MFLUW	Female	0	Yes	Yes	10
Yes						
19	4183-MYFRB	Female	0	No	No	21
Yes						
MultipleLines InternetService OnlineSecurity ... \						
0	No phone service		DSL		No	...
1		No	DSL		Yes	...
2		No	DSL		Yes	...
3	No phone service		DSL		Yes	...
4		No	Fiber optic		No	...
5		Yes	Fiber optic		No	...
6		Yes	Fiber optic		No	...
7	No phone service		DSL		Yes	...
8		Yes	Fiber optic		No	...

9	No	DSL	Yes	...
10	No	DSL	Yes	...
11	No	No internet service	...	
12	Yes	Fiber optic	No	...
13	Yes	Fiber optic	No	...
14	No	Fiber optic	Yes	...
15	Yes	Fiber optic	Yes	...
16	No	No internet service	...	
17	Yes	Fiber optic	Yes	...
18	No	DSL	No	...
19	No	Fiber optic	No	...
0	DeviceProtection	TechSupport	StreamingTV	\
1	No	No	No	
2	Yes	No	No	
3	No	No	No	
4	Yes	Yes	No	
5	No	No	Yes	
6	Yes	No	Yes	
7	No	No	No	
8	Yes	Yes	Yes	
9	No	No	No	
10	No	No	No	
11	No internet service	No internet service	No internet service	
12	Yes	No	Yes	
13	Yes	No	Yes	
14	Yes	Yes	Yes	
15	Yes	Yes	Yes	
16	No internet service	No internet service	No internet service	
17	Yes	No	Yes	
18	Yes	Yes	No	
19	Yes	No	No	
0	StreamingMovies	Contract	PaperlessBilling	\
1	No	Month-to-month	Yes	
2	No	One year	No	
3	No	Month-to-month	Yes	
4	No	One year	No	
5	No	Month-to-month	Yes	
6	Yes	Month-to-month	Yes	
7	No	Month-to-month	Yes	
8	No	Month-to-month	No	
9	Yes	Month-to-month	Yes	
10	No	One year	No	
11	No internet service	Two year	Yes	
12	Yes	One year	No	
13	Yes	Month-to-month	Yes	

```

14                 Yes Month-to-month          Yes
15                 Yes Two year             No
16 No internet service      One year           No
17                 Yes Two year             No
18                 No Month-to-month        No
19                 Yes Month-to-month        Yes

          PaymentMethod MonthlyCharges TotalCharges Churn
0       Electronic check      29.85        29.85   No
1       Mailed check         56.95     1889.50   No
2       Mailed check         53.85       108.15  Yes
3 Bank transfer (automatic) 42.30     1840.75   No
4       Electronic check      70.70       151.65  Yes
5       Electronic check      99.65       820.50  Yes
6 Credit card (automatic)  89.10     1949.40   No
7       Mailed check         29.75       301.90   No
8       Electronic check      104.80      3046.05 Yes
9 Bank transfer (automatic) 56.15     3487.95   No
10      Mailed check         49.95       587.45   No
11      Credit card (automatic) 18.95       326.80   No
12      Credit card (automatic) 100.35      5681.10   No
13 Bank transfer (automatic) 103.70      5036.30  Yes
14      Electronic check      105.50      2686.05   No
15      Credit card (automatic) 113.25      7895.15   No
16      Mailed check         20.65       1022.95   No
17 Bank transfer (automatic) 106.70      7382.25   No
18      Credit card (automatic) 55.20       528.35  Yes
19      Electronic check      90.05     1862.90   No

```

[20 rows x 21 columns]

```

# Drop customerID
df=df.drop('customerID',axis=1)
#axis : {0 or 'index', 1 or 'columns'}, default 0

# change target variable in binary format
df['Churn']=df['Churn'].map({'Yes':1,'No':0})

#df.to_excel("churn_cleaned_data_powerbi.xlsx", index=False)

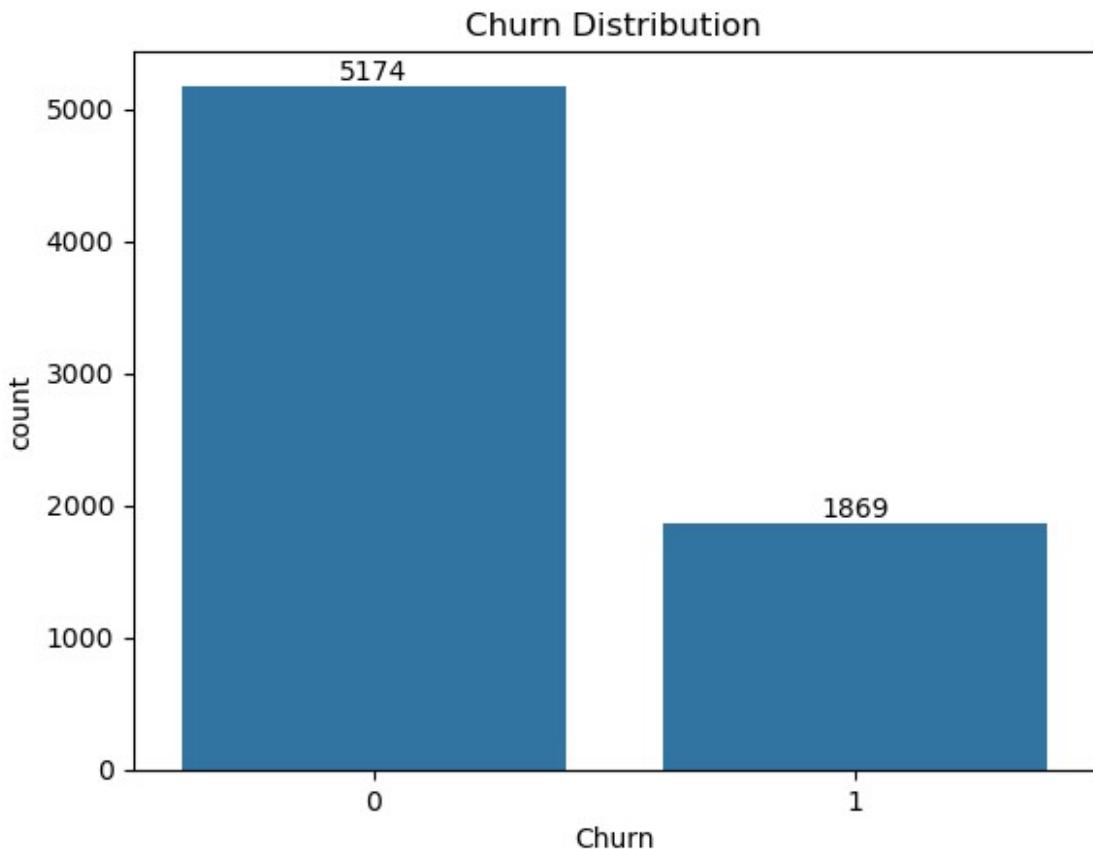
```

Exploratory Data Analysis (EDA)

```

# Check churn balance
ax=sns.countplot(x='Churn',data=df)
ax.bar_label(ax.containers[0])
plt.title('Churn Distribution')
plt.show()

```



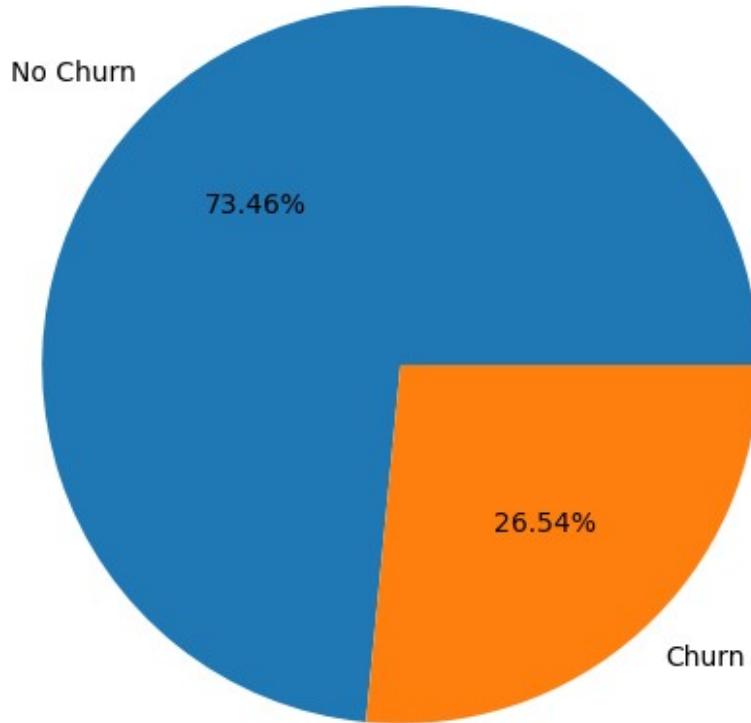
```
chrn_count=df['Churn'].value_counts()
chrn_count

Churn
0    5174
1    1869
Name: count, dtype: int64

plt.figure(figsize=(6,6))
plt.pie(chrn_count, labels=['No Churn', 'Churn'], autopct='%.1f%%')
plt.title('Percentage Of Churn Customer')

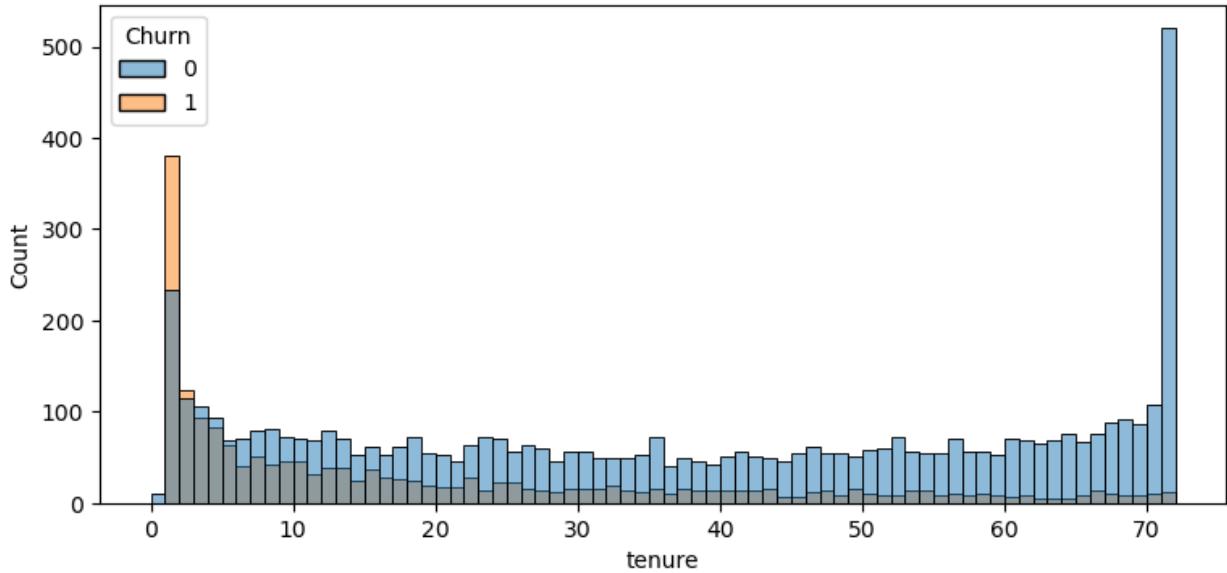
Text(0.5, 1.0, 'Percentage Of Churn Customer')
```

Percentage Of Churn Customer



in above we understand 73% people is not chrn and 27% people is chrun out

```
plt.figure(figsize=(9,4))
sns.histplot(x='tenure',data=df,bins=72,hue='Churn')
plt.show()
```



in above we understand who have used services is long time its stay who using only 1 and 2 month is churned

```

df.columns

Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
       'PhoneService', 'MultipleLines', 'InternetService',
       'OnlineSecurity',
       'OnlineBackup', 'DeviceProtection', 'TechSupport',
       'StreamingTV',
       'StreamingMovies', 'Contract', 'PaperlessBilling',
       'PaymentMethod',
       'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')

# Columns to plot
cols = [
    'PhoneService', 'MultipleLines', 'InternetService',
    'OnlineSecurity',
    'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
    'StreamingMovies'
]

# Create subplot grid (3 rows x 3 columns)
fig, axes = plt.subplots(3, 3, figsize=(15, 12))
axes = axes.flatten() # Flatten 2D array of axes for easy iteration

# Loop through columns
for i, col in enumerate(cols):
    sns.countplot(x=col, data=df, hue='Churn', ax=axes[i])
    axes[i].set_title(f'{col} vs Churn', fontsize=12)
    axes[i].bar_label(axes[i].containers[0], fontsize=8)

```

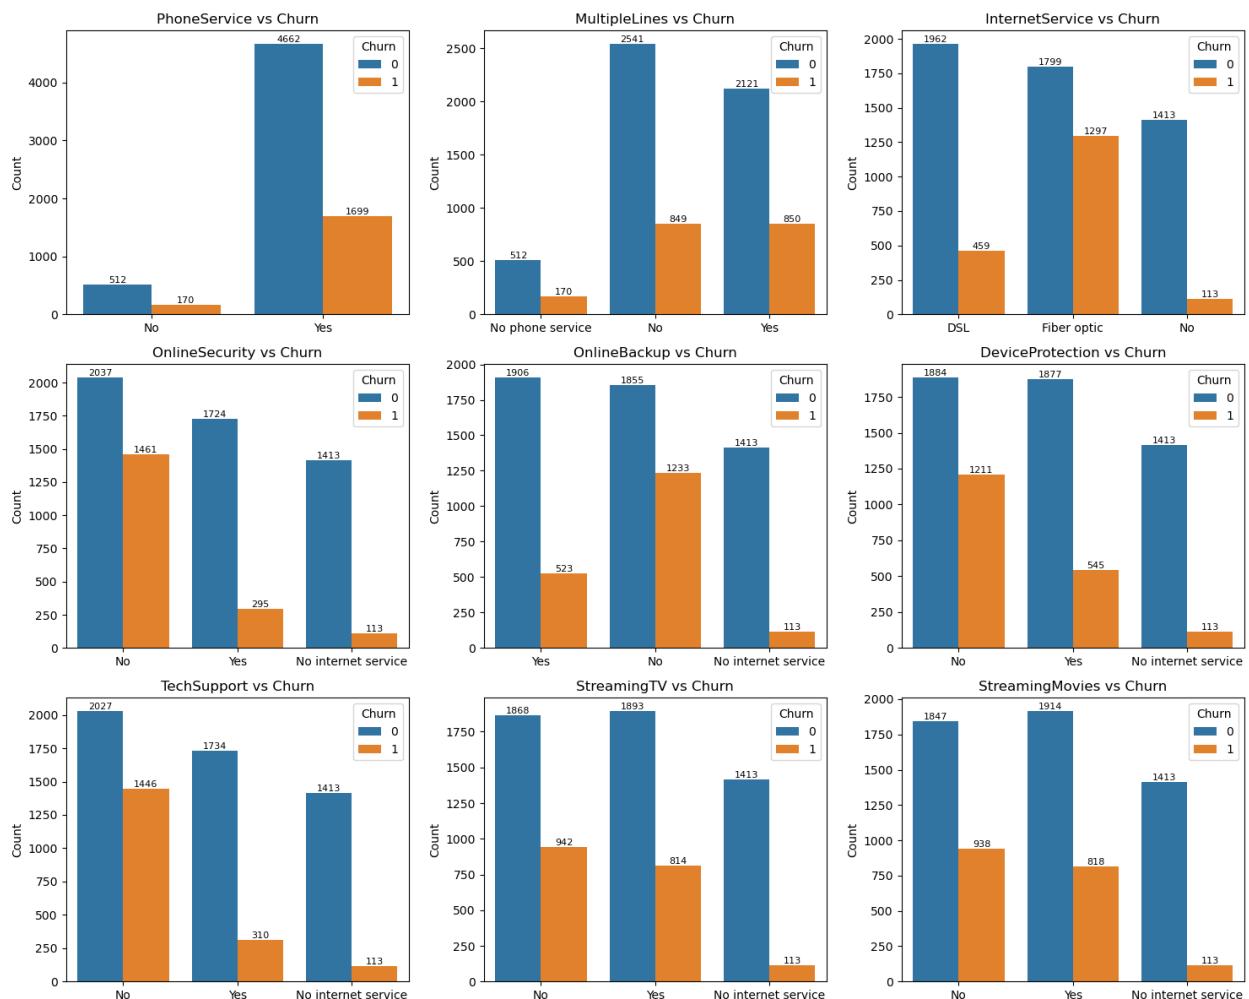
```

if len(axes[i].containers) > 1:
    axes[i].bar_label(axes[i].containers[1], fontsize=8)
axes[i].set_xlabel('')
axes[i].set_ylabel('Count')

# Remove extra empty plots (if any)
for j in range(len(cols), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```



in above Customers without internet-based add-ons (like OnlineSecurity, TechSupport, or DeviceProtection) tend to churn more frequently. Fiber optic users also exhibit higher churn rates compared to DSL users. Overall, having additional services (security, backup, or protection) is linked with lower churn

```

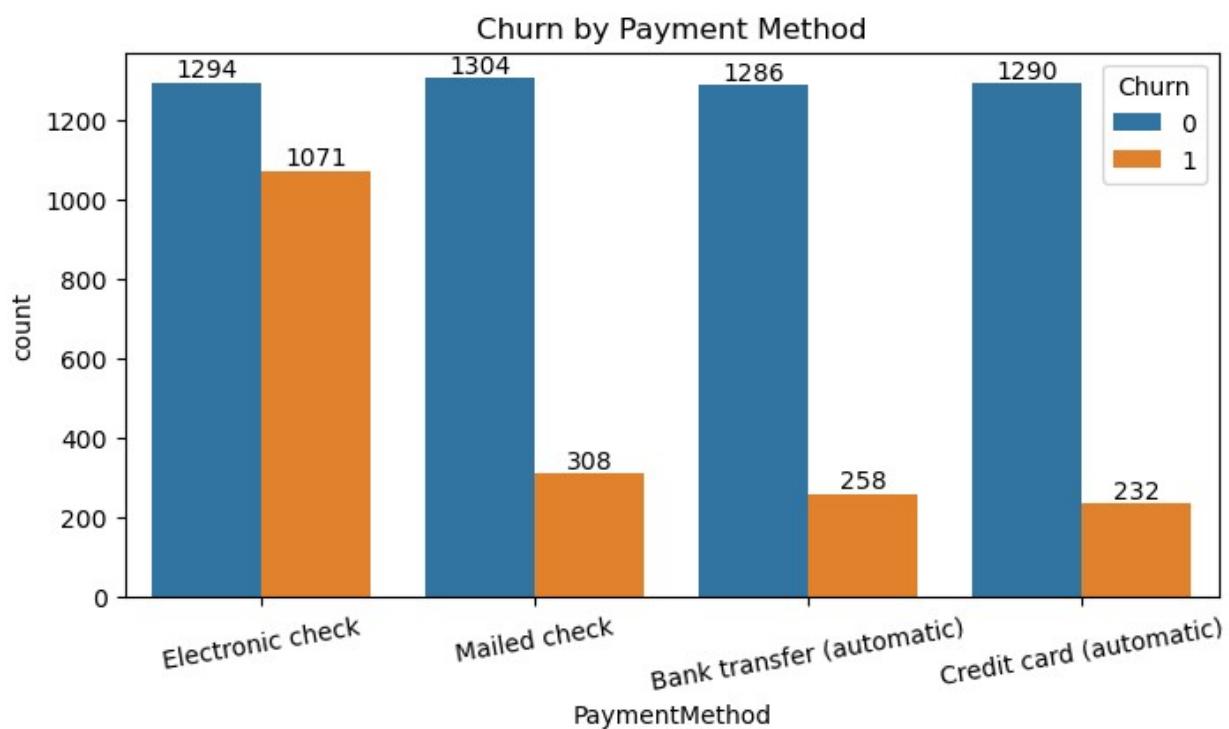
# analysis using paymnnet method
plt.figure(figsize=(8,4))

```

```

ax=sns.countplot(x='PaymentMethod',hue='Churn',data=df)
for containers in ax.containers:
    ax.bar_label(containers)
#ax.bar_label(ax.containers[0])
plt.title("Churn by Payment Method")
plt.xticks(rotation=10)
plt.show()

```



Customer is churn out when they are using Electronic Check

```

# Numeric summary
df.describe()

   SeniorCitizen      tenure  MonthlyCharges  TotalCharges
Churn
count    7043.000000  7043.000000    7043.000000  7043.000000
7043.000000
mean     0.162147    32.371149     64.761692   2281.916928
0.265370
std      0.368612    24.559481     30.090047   2265.270398
0.441561
min     0.000000    0.000000     18.250000   18.800000
0.000000
25%     0.000000    9.000000     35.500000   402.225000
0.000000
50%     0.000000   29.000000     70.350000  1397.475000
0.000000

```

```

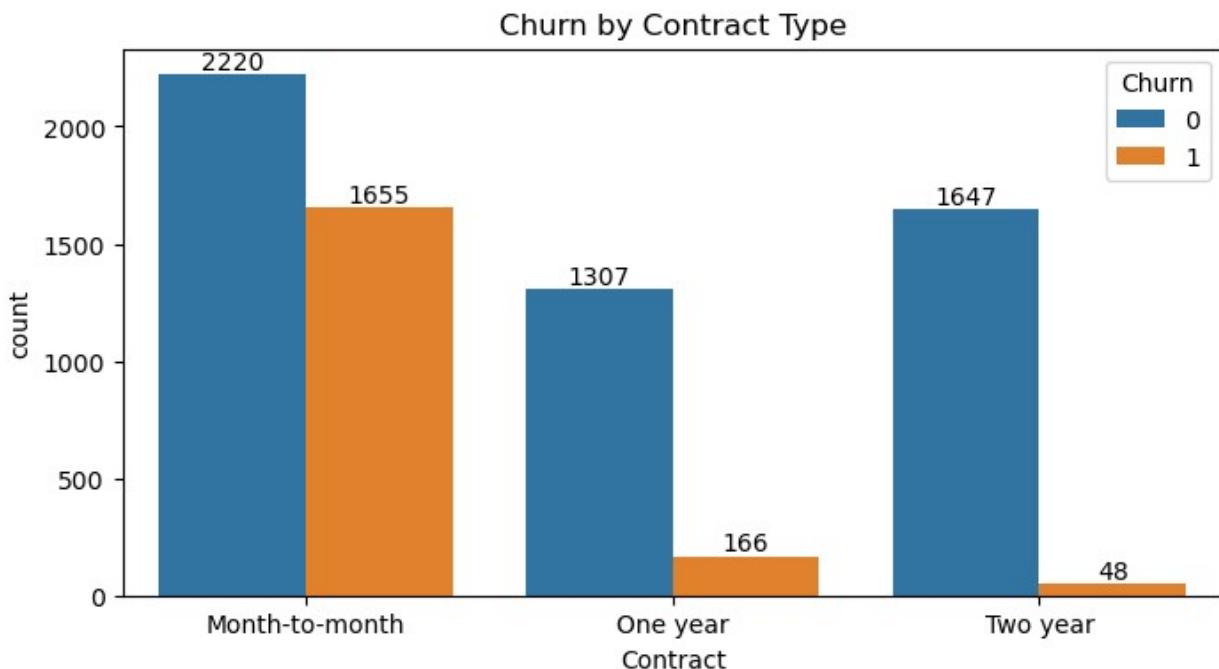
75%          0.000000    55.000000    89.850000  3786.600000
1.000000
max         1.000000    72.000000   118.750000  8684.800000
1.000000

```

```

# Categorical distribution example
plt.figure(figsize=(8,4))
ax=sns.countplot(x='Contract',hue='Churn',data=df)
for containers in ax.containers:
    ax.bar_label(containers)
#ax.bar_label(ax.containers[0])
plt.title("Churn by Contract Type")
plt.show()

```



in above you can say month to month contract base customer is churn more and for long term one year and 2 year customer is stay. for that we need to convey to customer contract for long time

Feature Encoding & Scaling

```

# Encode categorical columns(in this categorical columns store in 1 variable)
cat_col=df.select_dtypes(include='object').columns
cat_col

Index(['gender', 'Partner', 'Dependents', 'PhoneService',
       'MultipleLines',
       'InternetService', 'OnlineSecurity', 'OnlineBackup',
       'DeviceProtection',

```

```

'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
'PaperlessBilling', 'PaymentMethod'],
dtype='object')

#in this we convert categorical columns in True False Format and drop
first original column df =
df =pd.get_dummies(df,columns=cat_col,drop_first=True)
df

      SeniorCitizen  tenure  MonthlyCharges  TotalCharges  Churn
gender_Male \
0             0       1           29.85        29.85     0
False
1             0      34           56.95      1889.50     0
True
2             0       2           53.85        108.15     1
True
3             0      45           42.30      1840.75     0
True
4             0       2           70.70        151.65     1
False
...
...
7038            0      24           84.80      1990.50     0
True
7039            0      72           103.20      7362.90     0
False
7040            0      11           29.60        346.45     0
False
7041            1       4           74.40        306.60     1
True
7042            0      66           105.65      6844.50     0
True

      Partner_Yes  Dependents_Yes  PhoneService_Yes \
0          True           False           False
1         False           False           True
2         False           False           True
3         False           False          False
4         False           False           True
...
...
7038        True           True           True
7039        True           True           True
7040        True           True          False
7041        True           False           True
7042       False           False           True

      MultipleLines_No phone service  ...  StreamingTV_No internet
service \
0                  True   ...

```

```
False
1                               False ...
False
2                               False ...
False
3                               True ...
False
4                               False ...
False
...
...
7038                               False ...
False
7039                               False ...
False
7040                               True ...
False
7041                               False ...
False
7042                               False ...
False

StreamingTV_Yes  StreamingMovies_No internet service \
0                  False
1                  False
2                  False
3                  False
4                  False
...
7038                 True
7039                 True
7040                 False
7041                 False
7042                 True

StreamingMovies_Yes Contract_One year Contract_Two year \
0                  False
1                  False
2                  False
3                  False
4                  False
...
7038                 True
7039                 True
7040                 False
7041                 False
7042                 True

PaperlessBilling_Yes PaymentMethod_Credit card (automatic) \
0                  True
1                  False
```

```

1          False      False
2          True       False
3          False      False
4          True       False
...
7038        ...
7039        ...
7040        ...
7041        ...
7042        ...

```

	PaymentMethod_Electronic check	PaymentMethod_Mailed check
0	True	False
1	False	True
2	False	True
3	False	False
4	True	False
...
7038	False	True
7039	False	False
7040	True	False
7041	False	True
7042	False	False

[7043 rows x 31 columns]

```
# Split data in x and y
x = pd.get_dummies(df.drop('Churn', axis=1), drop_first=True)
feature_names = x.columns.tolist()
joblib.dump(feature_names, "feature_names.pkl")
y = df['Churn']

x.columns

Index(['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges',
       'gender_Male', 'Partner_Yes', 'Dependents_Yes',
       'PhoneService_Yes',
       'MultipleLines_No phone service', 'MultipleLines_Yes',
       'InternetService_Fiber optic', 'InternetService_No',
       'OnlineSecurity_No internet service', 'OnlineSecurity_Yes',
       'OnlineBackup_No internet service', 'OnlineBackup_Yes',
       'DeviceProtection_No internet service', 'DeviceProtection_Yes',
       'TechSupport_No internet service', 'TechSupport_Yes',
       'StreamingTV_No internet service', 'StreamingTV_Yes',
       'StreamingMovies_No internet service', 'StreamingMovies_Yes',
       'Contract_One year', 'Contract_Two year',
       'PaperlessBilling_Yes',
       'PaymentMethod_Credit card (automatic)',
       'PaymentMethod_Electronic check', 'PaymentMethod_Mailed
```

```

check'],
    dtype='object')

# Train-test split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42,stratify=y)
'''random_state : default=None
   Controls the shuffling applied to the data before applying the split.
shuffle : default=True
   Whether or not to shuffle the data before splitting. If shuffle=False
   then stratify must be None.'''
'random_state : default=None\n   Controls the shuffling applied to the data before applying the split.\nshuffle : default=True\nWhether or not to shuffle the data before splitting. If shuffle=False\n   then stratify must be None.'

# Scale
scaler=StandardScaler()
x_train=scaler.fit_transform(x_train)
x_test=scaler.transform(x_test)

```

Model Training (Logistic Regression & Random Forest)

```

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

models={'Logistic Regression':LogisticRegression(),
        'Random
Forest':RandomForestClassifier(n_estimators=100,random_state=42)}
'''n_estimators : int, default=100
   The number of trees in the forest.
random_state :default=None
   Controls both the randomness of the bootstrapping of the samples used
when building trees (if ``bootstrap=True``) and the sampling of the
features to consider when looking for the best split at each node'''

'n_estimators : int, default=100\n   The number of trees in the forest.\nrandom_state :default=None\n   Controls both the randomness of the bootstrapping of the samples used\nwhen building trees (if ``bootstrap=True``) and the sampling of the\nfeatures to consider when looking for the best split at each node'

for name,model in models.items():
    model.fit(x_train,y_train)

```

```

y_pred=model.predict(x_test)
print(f"== {name} ==")
print('accuracy score',accuracy_score(y_test,y_pred))
print("ROC-AUC:", roc_auc_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test,
y_pred))
print("-" * 40)

== Logistic Regression ==
accuracy score 0.8069552874378992
ROC-AUC: 0.7302823632746905
Classification Report:
precision    recall   f1-score   support
          0       0.85      0.89      0.87     1035
          1       0.66      0.57      0.61      374
          accuracy           0.81     1409
          macro avg       0.75      0.73      0.74     1409
          weighted avg     0.80      0.81      0.80     1409
          -----
          -----
== Random Forest ==
accuracy score 0.7856635911994322
ROC-AUC: 0.6918830246195975
Classification Report:
precision    recall   f1-score   support
          0       0.83      0.89      0.86     1035
          1       0.62      0.49      0.55      374
          accuracy           0.79     1409
          macro avg       0.73      0.69      0.70     1409
          weighted avg     0.77      0.79      0.78     1409
          -----
          -----

```

in above Logistic Regression outperforms Random Forest in this case, with higher accuracy, ROC-AUC, and recall. Therefore Logistic Regression It's the better choice for predicting churn in this dataset.

```

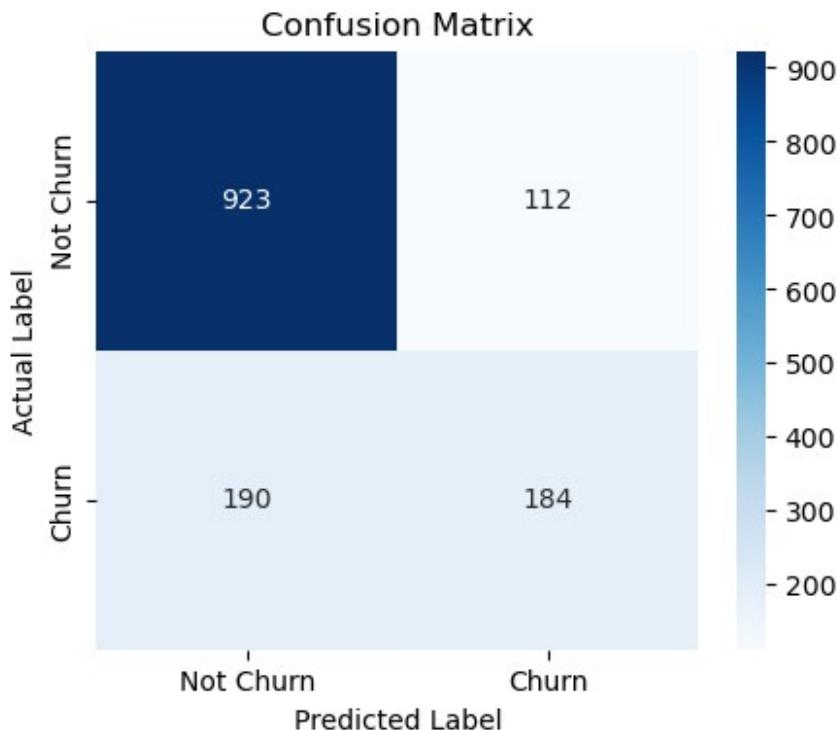
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Not Churn', 'Churn'],
            yticklabels=['Not Churn', 'Churn'])

plt.xlabel('Predicted Label')
plt.ylabel('Actual Label')

```

```
plt.title('Confusion Matrix')
plt.show()
```



Feature Importance (for Random Forest)

```
x.columns
Index(['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges',
       'gender_Male', 'Partner_Yes', 'Dependents_Yes',
       'PhoneService_Yes',
       'MultipleLines_No phone service', 'MultipleLines_Yes',
       'InternetService_Fiber optic', 'InternetService_No',
       'OnlineSecurity_No internet service', 'OnlineSecurity_Yes',
       'OnlineBackup_No internet service', 'OnlineBackup_Yes',
       'DeviceProtection_No internet service', 'DeviceProtection_Yes',
       'TechSupport_No internet service', 'TechSupport_Yes',
       'StreamingTV_No internet service', 'StreamingTV_Yes',
       'StreamingMovies_No internet service', 'StreamingMovies_Yes',
       'Contract_One year', 'Contract_Two year',
       'PaperlessBilling_Yes',
       'PaymentMethod_Credit card (automatic)',
       'PaymentMethod_Electronic check', 'PaymentMethod_Mailed
check'],
      dtype='object')
```

```

rf = RandomForestClassifier(random_state=42)
rf.fit(x, y)

RandomForestClassifier(random_state=42)

importances = pd.Series(rf.feature_importances_, index=x.columns)
importances.sort_values(ascending=False).head(10)

TotalCharges          0.192846
tenure                 0.173395
MonthlyCharges         0.172255
InternetService_Fiber optic  0.039823
PaymentMethod_Electronic check  0.035774
Contract_Two year      0.032939
gender_Male              0.028368
OnlineSecurity_Yes       0.026899
PaperlessBilling_Yes     0.025229
TechSupport_Yes           0.023075
dtype: float64

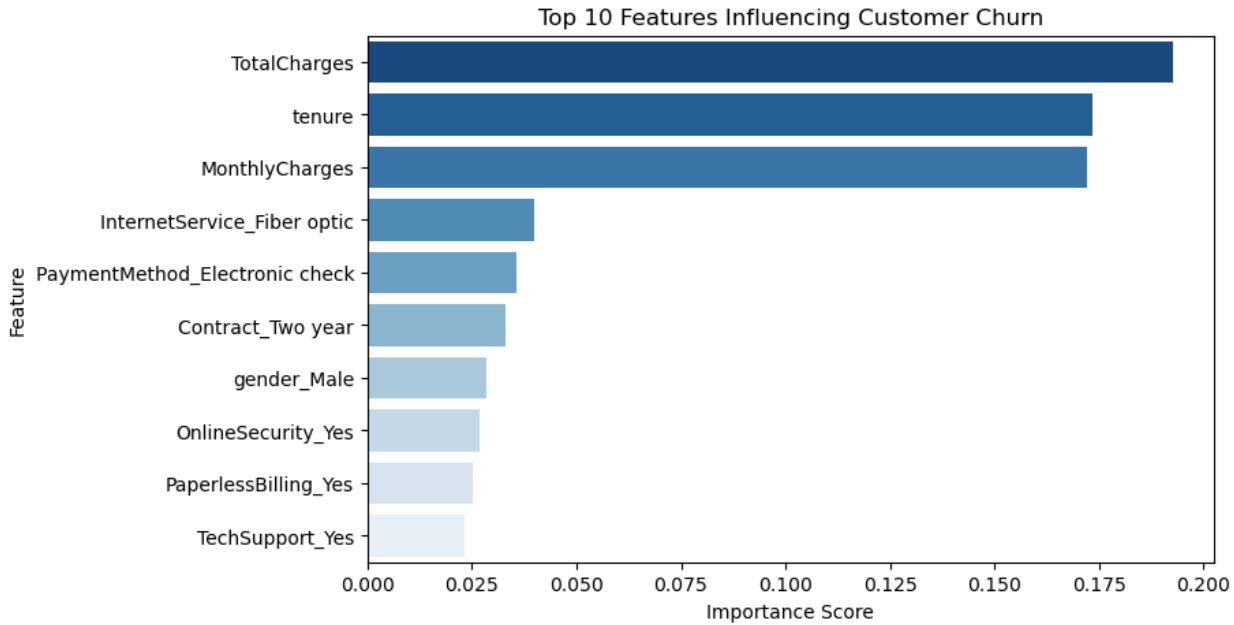
```

its tell you how important each feature was in making predictions.

```

plt.figure(figsize=(8,5))
sns.barplot(
    x=importances.nlargest(10),
    y=importances.nlargest(10).index,
    palette='Blues_r'
)
plt.title("Top 10 Features Influencing Customer Churn")
plt.xlabel("Importance Score")
plt.ylabel("Feature")
plt.show()

```



Financial factors (charges) and customer tenure are the strongest churn drivers. Service experience (security, support) helps retain customers. Payment and internet type also influence behavior, showing that both pricing and service quality matter. In short: High charges + short tenure + fiber optic users + electronic check payments = higher churn risk

improving recall for churn class

Use Class Weights

RandomForest supports class_weight='balanced' that forces model to pay more attention to minority class

```
rf = RandomForestClassifier(class_weight='balanced', random_state=42)
rf.fit(x_train, y_train)
y_pred = rf.predict(x_test)

print(classification_report(y_test, y_pred))

precision    recall   f1-score   support
          0       0.83      0.89      0.86     1035
          1       0.63      0.50      0.56      374
   accuracy                           0.79     1409
  macro avg       0.73      0.70      0.71     1409
weighted avg       0.78      0.79      0.78     1409
```

The use of class weights did not result in any significant improvement in recall.

Now We Are Using SMOTE Oversampling

```

from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state=42)

X_resampled, y_resampled = sm.fit_resample(x_train, y_train)

rf = RandomForestClassifier(random_state=42)
rf.fit(X_resampled, y_resampled)

y_pred = rf.predict(x_test)
print(classification_report(y_test, y_pred))

      precision    recall   f1-score   support
          0       0.85     0.84     0.85     1035
          1       0.57     0.60     0.59      374

   accuracy         0.78     1409
macro avg       0.71     0.72     0.72     1409
weighted avg    0.78     0.78     0.78     1409

```

Using SMOTE results are better than before — especially recall for churn (class 1) improving from 0.50 → 0.60. But we can still push churn recall higher

Now Try To Use Both SMOTE + Class Weight Together

```

sm = SMOTE(random_state=42)
x_train_res, y_train_res = sm.fit_resample(x_train, y_train)

rf = RandomForestClassifier(
    n_estimators=300,
    class_weight='balanced',
    random_state=42
)
rf.fit(x_train_res, y_train_res)

y_pred = rf.predict(x_test)
print(classification_report(y_test, y_pred))

      precision    recall   f1-score   support
          0       0.85     0.84     0.85     1035
          1       0.58     0.60     0.59      374

   accuracy         0.78     1409
macro avg       0.72     0.72     0.72     1409
weighted avg    0.78     0.78     0.78     1409

```

Using Both SMOTE + class_weight results are almost the same as pure SMOTE

Now we are using Threshold Tuning By default, prediction threshold = 0.50 So values like 0.40 or 0.35 are treated as NON-churn So we LOWER the threshold to 0.35 to catch more churn customers.

```
y_prob = rf.predict_proba(x_test)[:, 1]
```

in this only take the probability of class 1

```
y_pred_35 = (y_prob >= 0.35).astype(int)
print("Threshold = 0.35")
print(classification_report(y_test, y_pred_35))

Threshold = 0.35
      precision    recall  f1-score   support

          0       0.89     0.74     0.81     1035
          1       0.51     0.75     0.60      374

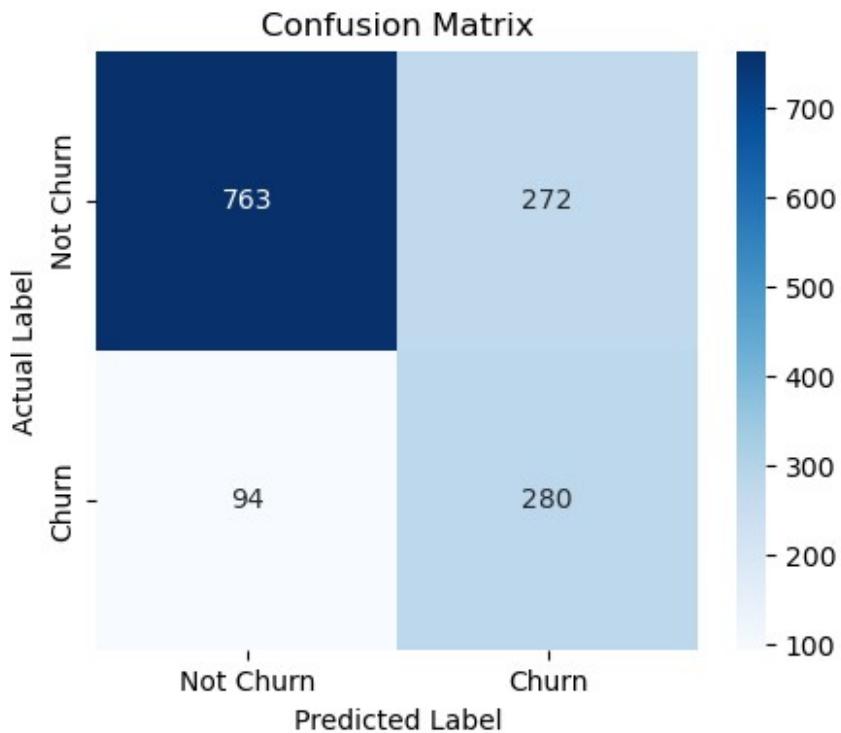
   accuracy                           0.74      1409
  macro avg       0.70     0.74     0.71      1409
weighted avg       0.79     0.74     0.75      1409
```

So Improved churn recall from 0.50 to 0.75 using probability threshold tuning (0.35)

```
cm = confusion_matrix(y_test, y_pred_35)

plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Not Churn', 'Churn'],
            yticklabels=['Not Churn', 'Churn'])

plt.xlabel('Predicted Label')
plt.ylabel('Actual Label')
plt.title('Confusion Matrix')
plt.show()
```



```
import joblib
joblib.dump(rf, "churn_model.pkl")
joblib.dump(scaler, "scaler.pkl")

['scaler.pkl']
```

```
import os
print(os.getcwd())

C:\Users\Admin
```